

Project Proposal: Explainable Machine Learning System for Real-Time Financial Fraud Detection

Team-3: Vinay kumar Simma (simmalv), Pradeep Yarlagadda (yarla3p), Praneeth Reddy Ambati (ambat3p)

Introduction

Financial fraud in online payment systems continues to grow in scale and complexity, creating significant challenges for automated detection systems. While modern machine learning models can identify suspicious transactions with high accuracy, many operate as "black boxes," offering limited insight into how decisions are made. This lack of transparency can reduce trust among analysts and slow down decision-making in high-risk financial environments.

This project proposes the development of an explainable, real-time machine learning framework for financial fraud detection using the Cifer-Fraud-Detection-Dataset-AF. By combining streaming data technologies with interpretable prediction models, the system aims to provide both accurate fraud classification and clear, human-readable explanations for each alert.

Primary Computer Science Areas of Focus

This project focuses on Machine Learning, Data Engineering, and Distributed Systems. The core technical emphasis is on supervised learning and anomaly detection models for fraud classification, combined with explainable AI (XAI) methods to interpret model behavior. The system will leverage a distributed architecture to handle high-frequency transaction streams.

Key Performance Indicators (KPIs)

- **Model Accuracy and Reliability:** Achieve high precision, recall, and F1-score in detecting fraudulent transactions.
- **Explanation Fidelity:** Ensure that XAI-generated explanations (SHAP/LIME) accurately represent the model's decision-making logic.
- **System Latency:** Maintain low-latency processing from transaction ingestion to alert generation with attached explanations.
- **Concept Drift Detection:** Successfully identify and report shifts in data distribution that impact model performance over time.

Data Sources

We will utilize the **Cifer-Fraud-Detection-Dataset-AF** (available via Hugging Face), a comprehensive 21-million-row synthetic dataset modeled after real-world financial logs.

Column Features: The dataset includes `step` (temporal unit), `type` (transaction category), `amount`, `oldbalanceOrig`, `newbalanceOrig`, `nameorig`, `namedest`, `isFraud`, `isFlaggedFraud`, `oldbalanceDest`, and `newbalanceDest`.

Target Labels: The `isFraud` column serves as the ground truth, while `isFlaggedFraud` provides a baseline for traditional rule-based detection.

Why this dataset?

Its accounting-style features (sender/receiver balance shifts) are ideal for SHAP and LIME explanations, as they allow the model to "explain" fraud through logical balance discrepancies.

Tech Stack

Machine Learning: Python, Scikit-learn, XGBoost, Random Forest, SHAP, and LIME for explainability.

Data Engineering & Streaming: Apache Kafka for simulating real-time transaction flows.

Backend & API: FastAPI for serving predictions and explanations.

Monitoring: Custom dashboard for visualizing concept drift and model performance metrics.

Project Description

The proposed system will follow a pipeline-oriented architecture:

1. **Data Ingestion:** A Kafka producer will stream records from the Cifer dataset to simulate live transactions.
2. **Preprocessing:** Real-time feature engineering to handle categorical encoding (`type`) and calculate balance deltas.
3. **Fraud Classification:** An ensemble of XGBoost and Random Forest models will classify incoming transactions.
4. **Explainability Layer:** For every flagged transaction, the system will generate a SHAP-based explanation highlighting the specific features (e.g., unusually high amount relative to balance) that triggered the alert.

5. **Drift Analysis:** The system will monitor for "concept drift," identifying when new fraud patterns emerge that differ from the training data.

Conclusion

By integrating the Cifer-Fraud-Detection-Dataset-AF with explainable AI techniques, this project will demonstrate how financial institutions can move beyond black-box models. The system enables human analysts to understand the reasoning behind automated fraud alerts, fostering trust and enabling more effective responses in high-stakes environments.

References:-

- [1] Dal Pozzolo, A., Bontempi, G., Snoeck, M., and Haenens, A. (2015). Adversarial Drift Detection in the Credit Card Fraud Detection Domain. *Intelligent Systems in Accounting, Finance and Management*, 24(1), pp.37–48.
- [2] Carcillo, F., Dal Pozzolo, A., Bontempi, G., and Snoeck, M. (2021). Scarff: A Scalable Framework for Streaming Credit Card Fraud Detection with Drift Awareness. *Information Fusion*, 72, pp.182–194.
- [3] Ribeiro, M. T., Singh, S., and Guestrin, C. (2016). "Why Should I Trust You?" Explaining the Predictions of Any Classifier. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*.
- [4] Lundberg, S. M. and Lee, S.-I. (2017). A Unified Approach to Interpreting Model Predictions. *Advances in Neural Information Processing Systems*, 30.
- [5] Bahnsen, A. C., Aouada, D., and Ottersten, B. (2015). Cost-Sensitive Decision Trees for Fraud Detection. *Expert Systems with Applications*, 42(7), pp.3769–3783.